

Constraint-Aware Variable Ordering in Symbolic Optimal Planning

Final Course Project Report – Artificial Intelligence and Autonomous Systems (00960208)

Mozes Yaron

Faculty of Data and Decision Sciences
Technion – Israel Institute of Technology
yaron.mozes@campus.technion.ac.il

Kadzelshvily Galit

Faculty of Data and Decision Sciences
Technion – Israel Institute of Technology
galit.k@campus.technion.ac.il

August 2026 · <https://github.com/YaronMozes/slbd-symbolic-search-research>

Abstract

Symbolic BDD-based search is a prominent paradigm for cost-optimal classical planning, with its performance heavily dictated by BDD sizes and, consequently, variable ordering algorithms. State-of-the-art planners such as SymK rely on causal-graph proximity heuristics (e.g., GAMER) but ignore the structural impact of the mutex and invariant constraints compiled into the objects search manipulates. We introduce constraint-aware variable ordering - incorporating mutex and invariant co-occurrence edges into the GAMER-family linear arrangement objective - and evaluate its impact across the complete IPC optimal-STRIPS suite within SymK (IPC-2023). We show that a targeted per-domain policy gains coverage over stock SymK on the five domains it targets, reproducibly across replicates (though those domains were selected post hoc), and yields substantial search acceleration in domains such as *woodworking*. Suite-wide, however, enabling the ordering universally results in a tightly bounded null effect. Instrumentation explains the discrepancy: while the ordering successfully shrinks transition relations, transition relation size is largely decoupled from search cost, which is instead dominated by peak search-BDD size. Finally, we highlight key methodological pitfalls, including regression to the mean in "hard instance" subsets and cross-year IPC duplicate leakage in domain-wise cross-validation.

1. Introduction

Cost-optimal classical planning aims to find provably minimal-cost plan sequences in deterministic state spaces. A dominant paradigm for blind optimal planning is symbolic search, which represents sets of reachable states and transition relations compactly as Binary Decision Diagrams (BDDs) [10]. Because symbolic search operations scale with the physical node counts of the underlying BDDs rather than the raw number of states, search efficiency is strictly governed by BDD size [5]. BDD sizes, in turn, depend critically (and often exponentially) on the choice of state variable ordering [5]. Modern state-of-the-art symbolic planners, including SymK [7, 6] – which

serves as the symbolic search component of Ragnarok, the winner of the IPC-2023 optimal track [2, 8] – inherit variable ordering algorithms directly from the GAMER family [3]. These algorithms compute variable orders by minimizing a weighted linear arrangement objective over edges of causal-graph proximity [3].

In addition to causal transitions, modern symbolic planners heavily rely on state-invariant constraints [9, 11]. In particular, h^2 mutexes and exactly-one invariant groups are compiled into constraint BDDs used to remove spurious states [9, 11]. Under the default *e-deletion* scheme these constraints are pushed into each operator's transition relation and applied to the frontier at initialization, so they shape the objects every image operation touches; under MUTEX_AND they are instead conjoined at every expansion. In either case the resulting BDD sizes depend on the variable ordering, yet standard causal-graph ordering objectives ignore the constraints entirely [9]. Over a decade ago, Torralba and Alcázar [9] noted this gap and suggested investigating variable orders derived from the constraints as future work. Our method is one concrete realisation of that broad direction: mutex and exactly-one co-occurrence edges inside the GAMER arrangement objective. To our knowledge the suggestion has not been implemented or evaluated in a state-of-the-art planner.

Core Research Question

Can conjoining mutex and invariant co-occurrence edges into the GAMER variable ordering objective systematically compress constraint BDDs and accelerate state-of-the-art symbolic search in SymK?

Variable ordering was not our first target: we began with search-control alternatives – landmark-guided direction selection for bidirectional symbolic search, meet-in-the-middle BDD-overlap signals, and forward-reachability pruning of the backward frontier – and found each to be net-neutral or actively harmful, because none of them addresses the structural BDD-size bottleneck. In this work, we therefore implement and systematically evaluate constraint-aware variable ordering inside SymK

on the complete IPC optimal-STRIPS benchmark suite (66 domains, 1,847 instances). We also correct a long-standing defect in the deployed GAMER ordering optimizer, which silently ignores edge weights during local search swaps.

We show that a targeted, pre-registered per-domain policy gains +6 coverage in each of two replicates on its five (post-hoc selected) target domains while accelerating domains like *woodworking* by 1.55 $\times$. Enabled universally, however, the ordering produces a tightly bounded null effect because it compresses transition relations without shrinking peak search-BDD sizes, the true runtime bottleneck.

2. Problem Formulation

2.1 SAS⁺ Planning and Invariants

A classical planning task in SAS⁺ representation is defined as a tuple $\Pi = \langle \mathcal{V}, \mathcal{A}, s_0, s_g, \text{cost} \rangle$ [5]:

- $\mathcal{V} = \{v_1, \dots, v_n\}$ is a finite set of state variables. Each variable $v \in \mathcal{V}$ has a finite domain $\mathcal{D}(v)$. A *fact* is a pair (v, d) with $d \in \mathcal{D}(v)$.
- A *state* s is a complete assignment to $\mathcal{V}$, where $s(v) \in \mathcal{D}(v)$ for all $v \in \mathcal{V}$. The state space is denoted $\mathcal{S} = \prod_{v \in \mathcal{V}} \mathcal{D}(v)$.
- $s_0 \in \mathcal{S}$ is the initial state, and s_g is a partial state assignment representing the goal condition.
- $\mathcal{A}$ is a finite set of actions. Each action $a \in \mathcal{A}$ is a pair $a = \langle \text{pre}(a), \text{eff}(a) \rangle$, where $\text{pre}(a)$ and $\text{eff}(a)$ are partial assignments over $\mathcal{V}$. Action execution incurs a non-negative cost $\text{cost}(a) \in \mathbb{R}_{\geq 0}$.

An action a is applicable in s if $s \models \text{pre}(a)$. Applying a in s yields state $s' = a(s)$ where $s'(v) = \text{eff}(a)[v]$ for variables defined in $\text{eff}(a)$, and $s'(v) = s(v)$ otherwise. A *cost-optimal plan* is a sequence of actions $\pi = \langle a_1, \dots, a_k \rangle$ transforming s_0 into a goal state $s \models s_g$ while minimizing $\sum_{i=1}^k \text{cost}(a_i)$.

In addition to causal dynamics, domain abstractions yield state invariants I :

1. **h^2 Mutexes** [9]: Pairs of unreachable fact combinations $\mu = \{(v_i, d_i), (v_j, d_j)\}$ such that no reachable state $s \in \mathcal{S}$ satisfies $s(v_i) = d_i \wedge s(v_j) = d_j$.
2. **Exactly-One Invariants** [9]: Subsets of facts $\mathcal{G} \subset \bigcup_v \{(v, d) \mid d \in \mathcal{D}(v)\}$ such that every reachable state satisfies $\sum_{f \in \mathcal{G}} \mathbb{I}(s \models f) = 1$.

2.2 Symbolic Search and Constraint Conjunction

Symbolic planning algorithms represent sets of states $S \subseteq \mathcal{S}$ and transition relations $T \subseteq \mathcal{S} \times \mathcal{S}$ as Reduced Ordered Binary Decision Diagrams (BDDs) [5, 10].

To map SAS⁺ variables to Boolean BDD variables, each $v_i \in \mathcal{V}$ is encoded using $m_i = \lceil \log_2 |\mathcal{D}(v_i)| \rceil$ bits. Let

$X = \{x_1, \dots, x_M\}$ denote the complete set of unprimed Boolean variables representing current states, where $M = \sum_{i=1}^n m_i$, and let $X' = \{x'_1, \dots, x'_M\}$ denote primed variables representing successor states.

Definition (Variable Ordering)

A **variable ordering** π is a bijection $\pi : X \cup X' \rightarrow \{1, \dots, 2M\}$ specifying the linear topological sequence in which Boolean variables are evaluated along all root-to-leaf paths in the BDD.

In practice the search is over a coarser space: the optimizer permutes the n SAS⁺ variables, and each variable's m_i bits – interleaved with their primed counterparts – occupy a contiguous block of the induced Boolean order. All objectives below are therefore stated over SAS⁺ variable positions.

In symbolic uniform-cost search (e.g., *sym-bd* in SymK) [7, 10], forward search expands a frontier BDD $F_k(X)$ at cost level k through a *relational product* (image computation). Writing the constraint restriction explicitly, the reachable successor set at level $k+1$ is [9]:

$$F_{k+1}(X') = \exists X (F_k(X) \wedge T(X, X') \wedge C(X')) \quad (1)$$

where $T(X, X') = \bigvee_{a \in \mathcal{A}} T_a(X, X')$ represents the system transition relation, and $C(X')$ is the invariant constraint characterising valid states, compiling all h^2 mutexes and exactly-one groups [9]:

$$C(X) = \left(\bigwedge_{\mu \in M_{h^2}} \neg \mu(X) \right) \wedge \left(\bigwedge_{\mathcal{G} \in I_{\text{one}}} \text{ExactlyOne}(\mathcal{G}, X) \right) \quad (2)$$

Equation 1 states the semantics; the schedule differs by configuration. Our experiments use SymK's default *e-deletion*, in which C is compiled into the transition relations T_a ahead of search and applied to the frontier at initialization rather than re-conjoined at each step. The variable ordering therefore governs the size of both T and C , and BDD widths are bounded by that ordering [5] – a dependency that constraint-graph ordering heuristics such as FORCE exploit outside of planning [1].

2.3 Causal Variable Ordering Objectives (GAMER)

State-of-the-art planners such as SymK [7] compute variable orderings using the GAMER local search framework [3, 5]. GAMER constructs an influence graph $G_{\text{causal}} = (\mathcal{V}, E_{\text{causal}}, w_{\text{causal}})$, where edges $(v_i, v_j) \in E_{\text{causal}}$ correspond to causal dependencies (e.g., co-occurrence in action preconditions or effects).

GAMER determines an optimal permutation $\pi^* \in \mathcal{P}(\mathcal{V})$ by minimizing a weighted linear arrangement (WLA) energy objective [3, 5]:

$$\mathcal{J}_{\text{causal}}(\pi) = \sum_{(v_i, v_j) \in E_{\text{causal}}} w_{\text{causal}}(v_i, v_j) |\pi(v_i) - \pi(v_j)|^2 \quad (3)$$

2.4 Problem Statement: Constraint-Aware Variable Ordering

While Equation 3 optimizes causal variable proximity to shrink transition relations $T_a(X, X')$, it completely ignores the structural dependencies of the constraint BDD $C(X')$ in Equation 2 [9]. Because C is compiled into the transition relations that every image operation applies, a variable ordering that is poor for C propagates into the objects search manipulates throughout [9].

Our goal is to formalize a generalized objective function $\mathcal{J}_{\text{combined}}(\pi)$ that incorporates constraint co-occurrence structures $E_{\text{constraint}}$ directly into the linear arrangement search:

$$\mathcal{J}_{\text{combined}}(\pi) = \mathcal{J}_{\text{causal}}(\pi) + \lambda \cdot \mathcal{J}_{\text{constraint}}(\pi) \quad (4)$$

where $\lambda \in \mathbb{R}_{\geq 0}$ parameterizes the relative weight of invariant constraints relative to causal transition dynamics. In our implementation this trade-off parameter is realized directly as the base constraint edge weight w_{co} introduced in Section 3.1, i.e. $\lambda \equiv w_{\text{co}}$ with causal edges normalized to unit weight.

3. Methodology

3.1 Constraint-Aware Influence Graph Construction

Classical GAMER ordering constructs an undirected influence graph $G = (\mathcal{V}, E, w)$ over SAS⁺ variables $\mathcal{V}$ where E consists solely of causal dependencies E_{causal} derived from action preconditions and effects [3, 5]. To optimize for the constraint structure (C in Equation 2) that e-deletion compiles into the transition relations and initial frontier, we expand G by overlaying a set of co-occurrence edges $E_{\text{constraint}}$ derived from h^2 mutexes and exactly-one invariant groups [9].

Let $\mathcal{M}$ denote the set of all static invariant groups. Each group $M \in \mathcal{M}$ contains a set of facts $\mathcal{F}_M = \{(v_1, d_1), \dots, (v_k, d_k)\}$. For every pair of distinct facts $(v_i, d_i), (v_j, d_j) \in \mathcal{F}_M$ with $v_i \neq v_j$, we insert an edge (v_i, v_j) into $E_{\text{constraint}}$ with a base weight $w_{\text{co}} \in \mathbb{R}_{>0}$.

Sections 3.1.1 and 3.1.2 describe two optional refinements of this construction. Both are exposed as ablation switches and, as reported in Section 3.4, *neither is enabled in the pre-registered confirmed policy*; we describe them because they delimit the design space we explored.

3.1.1 Group-Size Normalization (optional)

An invariant group containing $k = |\mathcal{F}_M|$ facts generates $\binom{k}{2} = \frac{k(k-1)}{2}$ variable pairs. Unnormalized, large invariant groups (e.g., domain location features where $k \geq 30$) produce $\mathcal{O}(k^2)$ edge pairs that heavily swamp the sparse causal graph E_{causal} . To prevent dense invariant groups from drowning out causal structure, we define a normalized edge weight w_{pair} :

$$w_{\text{pair}}(M) = \frac{w_{\text{co}}}{\max(1, |\mathcal{F}_M| - 1)} \quad (5)$$

This caps the total accumulated influence weight contributed by any single invariant group M to $\approx \frac{w_{\text{co}} \cdot k}{2}$, maintaining a balanced weight ratio between pairwise h^2 mutexes ($k = 2$, weight w_{co}) and large invariant groups.

3.1.2 Causal Edge Skipping (optional)

In domains with dense causal topologies (e.g., *blocks-world*, where any block can interact with any other), almost every constraint pair $(v_i, v_j) \in E_{\text{constraint}}$ is already connected in E_{causal} . Over-constraining these pairs tightens the variable ordering excessively, restricting local search freedom and degrading search performance. We introduce an optional filtering predicate $\text{SkipCausal}(v_i, v_j)$:

$$\text{SkipCausal}(v_i, v_j) = \mathbb{I}((v_i, v_j) \in E_{\text{causal}} \vee (v_j, v_i) \in E_{\text{causal}})$$

When enabled, constraint edges that duplicate existing causal graph adjacencies are discarded, preserving long-range constraint edges that assist sparse domains (e.g., *depot*).

3.1.3 Combined Weight Accumulation

The total weight $w(v_i, v_j)$ between any two state variables $v_i, v_j \in \mathcal{V}$ in the combined influence graph is computed as:

$$w(v_i, v_j) = w_{\text{causal}}(v_i, v_j) + \sum_{M \in \mathcal{M}_{i,j}} w_{\text{pair}}(M) \cdot (1 - \text{SkipCausal}(v_i, v_j)) \quad (6)$$

where $\mathcal{M}_{i,j}$ is the set of invariant groups containing co-occurring facts for variables v_i and v_j , and each group contributes once *per co-occurring fact pair* (a group holding several facts of v_i and v_j adds w_{pair} with the corresponding multiplicity). In the confirmed policy configuration this reduces to $w_{\text{pair}}(M) = w_{\text{co}}$ and $\text{SkipCausal} \equiv 0$.




Figure 1: System pipeline for Constraint-Aware Variable Ordering in SymK.

3.2 The Weight-Aware Local Search Optimizer Repair

Non-uniform edge weighting (Equation 6) requires the local search optimizer to evaluate edge weights faithfully during permutation steps. The optimizer deployed in SymK does not.

3.2.1 The Deployed Defect

In stock GAMER (inherited by SymK), the local search arrangement function evaluated edge *existence* rather than edge *weight*. Specifically, during the $\mathcal{O}(|\mathcal{V}|)$ incremental swap delta calculation when swapping positions of variables at index p_1 and p_2 , the code checked ‘if (influence(v_i, v_j))’ as a boolean existence flag. As a result, non-unit weights $w(v_i, v_j) \neq 1.0$ were silently flattened to 1.0. Any weighted topology blend was transformed into an unweighted graph dominated by the raw pair counts of large mutex groups.

3.2.2 The Weight-Aware Objective Repair

We repaired the objective function and the incremental swap delta evaluation loop. Let $\pi(p)$ denote the variable placed at topological position $p \in \{1, \dots, |\mathcal{V}|\}$. The total linear arrangement cost $\mathcal{J}(\pi)$ is defined as:

$$\mathcal{J}(\pi) = \sum_{1 \leq p_a < p_b \leq |\mathcal{V}|} w(\pi(p_a), \pi(p_b)) \cdot (p_b - p_a)^2 \quad (7)$$

When proposing a candidate swap between positions p_1 and p_2 (with $p_1 \neq p_2$), write $w_{k,i} := w(\pi(k), \pi(p_i))$ for the weight between the variable at position k and the variable at swap position p_i . The change in objective value $\Delta\mathcal{J}(p_1, p_2)$ is then computed incrementally in $\mathcal{O}(|\mathcal{V}|)$ time:

$$\Delta\mathcal{J}(p_1, p_2) = \sum_{k \neq p_1, p_2} \left[w_{k,1}((k - p_2)^2 - (k - p_1)^2) + w_{k,2}((k - p_1)^2 - (k - p_2)^2) \right] \quad (8)$$

If $\Delta\mathcal{J}(p_1, p_2) < 0$, the swap is accepted.

Backward Compatibility Guarantee

For stock GAMER causal graphs, all edge weights are identically $w(v_i, v_j) = 1.0$. Consequently, Equation 7 and Equation 8 yield bit-identical variable orderings to the original unweighted baseline while correctly honoring weighted topologies.

3.3 Algorithmic Implementation

Algorithm 1 details the complete variable ordering pipeline.

3.4 Pre-Registered Per-Domain Policy Execution

Because constraint density varies by domain structure, applying constraint ordering suite-wide induces domain-specific trade-offs. To evaluate the replicated per-domain signals under a fixed protocol, we define a targeted per-domain policy $\mathcal{W}(d)$ mapping each domain d to its con-

Algorithm 1: Constraint-Aware Variable Ordering Engine

Input: SAS⁺ Task II, Mutex Groups $\mathcal{M}$, Base Weight w_{co} , Iters N_{iter} , Restarts R

Output: Optimized Variable Order π^*

1. Initialize $G = (\mathcal{V}, E, w)$ with $w(u, v) \leftarrow 0$ for all $u, v \in \mathcal{V}$
2. **if** $\neg \text{ConstraintOnly}$ **then**
3. **for each** causal edge $(u, v) \in E_{\text{causal}}(\Pi)$ **do**
 $w(u, v) \leftarrow 1.0$
4. **end if**
5. **for each** mutex group $M \in \mathcal{M}$ **do**
6. $w_{\text{pair}} \leftarrow \text{Normalize}(w_{\text{co}}, |M|)$ via Eq. 5 [opt.]
7. **for each** pair $(v_i, d_i), (v_j, d_j) \in \text{Facts}(M)$ with $v_i \neq v_j$ **do**
8. **if** $\neg \text{SkipCausal}(v_i, v_j)$ **then** [opt.]
9. $w(v_i, v_j) \leftarrow w(v_i, v_j) + w_{\text{pair}}$
10. **end if**
11. **end for**
12. **end for**
13. $\pi^* \leftarrow \text{IdentityOrder}(\mathcal{V}); \quad \mathcal{J}^* \leftarrow \mathcal{J}(\pi^*)$ via Eq. 7
14. **for** $r = 0 \dots R$ **do**
15. $\pi \leftarrow (r == 0) ? \pi^* : \text{RandomPermutation}(\mathcal{V})$
16. **for** $i = 1 \dots N_{\text{iter}}$ **do**
17. Sample distinct positions $p_1, p_2 \sim \text{Uniform}(1, |\mathcal{V}|)$
18. Compute $\Delta\mathcal{J}(p_1, p_2)$ via Eq. 8
19. **if** $\Delta\mathcal{J}(p_1, p_2) < 0$ **then** $\text{Swap}(\pi[p_1], \pi[p_2])$
20. **end for**
21. **if** $\mathcal{J}(\pi) < \mathcal{J}^*$ **then** $\pi^* \leftarrow \pi; \quad \mathcal{J}^* \leftarrow \mathcal{J}(\pi)$
22. **end for**
23. **return** π^*

[opt.] Steps 6 and 8 are optional ablation switches. The pre-registered confirmed policy (Section 3.4) uses **neither**: $w_{\text{pair}} = w_{\text{co}}$ and $\text{SkipCausal} \equiv 0$. Configuration: $N_{\text{iter}} = 50,000$, $R = 20$, fixed RNG seed.

straint edge weight:

$$\mathcal{W}(d) = \begin{cases} 2.0 & \text{if } d \in \{\text{woodworking-opt08}, \\ & \text{woodworking-opt11}\} \\ 1.0 & \text{if } d \in \{\text{parking-opt14}, \\ & \text{barman-opt11, tpp}\} \\ 0.0 & \text{otherwise} \end{cases} \quad (9)$$

When $\mathcal{W}(d) = 0.0$, the constraint accumulation loop is bypassed completely, reverting to stock SymK by construction. For exact reproduction: the confirmed policy sets $w_{\text{co}} = \mathcal{W}(d)$ and applies it *without* group-size normalization (Section 3.1.1) and *without* causal-edge skipping (Section 3.1.2).

4. Results

We evaluate our constraint-aware variable ordering framework inside SymK across the complete International Planning Competition (IPC) optimal-STRIPS benchmark suite, comprising 66 domains and 1,847 instances. All experiments were conducted under standardized resource limits of 1800 seconds wall-clock time, 8 GB RAM, and 1 CPU core per task. The design and pass/fail gates of our headline evaluation were committed publicly before the confirmatory results were produced (pre-registration document in the repository).

4.1 Per-Domain Gains Under a Pre-Registered A/B

Initial full-suite exploratory runs identified strong, replicated positive signals on five specific domains: *parking-opt14*, *barman-opt11*, *tpp*, *woodworking-opt08*, and *woodworking-opt11*. On *woodworking*, varying the constraint weight w_{co} exhibited a monotonic dose-response acceleration ($w = 0.5 \rightarrow 0.87\times$, $w = 1.0 \rightarrow 0.84\times$, $w = 2.0 \rightarrow 0.586\times$ relative to stock, faster on 38/47 instances, exact Wilcoxon $p = 4 \cdot 10^{-7}$, sign test $p = 2.5 \cdot 10^{-5}$). This exploratory figure is a search-time ratio; the confirmatory study below measures the same effect end-to-end in wall-clock at 0.647 (1.55 $\times$), and we quote the latter throughout.

To validate this signal, we pre-registered a targeted per-domain policy: enabling constraint-aware variable ordering with per-domain weights $\mathcal{W}(d)$ (Equation 9) on the five signal domains while executing stock SymK elsewhere. We executed a confirmatory, double-replicated A/B study with 480 runs, placing each instance’s stock and policy jobs adjacently in the execution queue so that the two arms meet similar machine load. Adjacency reduces but does not eliminate order effects: the queue always emits stock before policy, and we did not counter-balance.

As summarized in Table 1, **all four pre-registered evaluation gates passed**. The per-domain coverage gains relative to stock SymK across both execution replicates are detailed in Table 2.

Headline Result: Repeated +6 Coverage on the Five Target Domains

On the five domains it targets, the pre-registered policy gains +6 coverage in each of two replicates (64/65 $\rightarrow$ 70/71 of 120 instances) and accelerates *woodworking* by 1.55 $\times$, with no cost disagreement on any commonly solved instance (domain selection and scope: Section 5.4).

4.2 Bounded Suite-Wide Null for Always-On Ordering

When constraint-aware variable ordering is enabled universally across all 66 domains, the suite-wide performance changes minimally. Across 1,847 benchmark instances, stock SymK solves 1145 instances, whereas the always-on constraint ordering solves 1139 instances.

The overall coverage difference (−6 instances) is statistically unremarkable: discordant pairs show 13 gains versus 19 losses (exact sign test $p = 0.38$), and we have no clean full-suite replicate from which to calibrate how much coverage drift this protocol produces on its own (Section 4.4).

On commonly solved instances ($n = 1,126$), the wall-clock geometric mean ratio is 1.005 with a median of 1.000 and a domain-clustered 95% confidence interval of [0.97, 1.05] (Table 3). This bounds the always-on ordering’s mean wall-clock effect on commonly solved instances to within about 5%; it does not bound coverage, which is addressed above.

Seed control evaluations demonstrate that the ordering induces real per-instance perturbations ($\text{sd}(\log \text{ wall ratio}) = 0.43$ vs. 0.24 for identical-code-path reruns), but the net distribution is unbiased. Isolated suite-wide regressions occur on domains like *freecell* (coverage drops 25 $\rightarrow$ 20, replicated) and *floortile-opt14* (1.56 $\times$ median, slower on 20/20 commonly solved instances, sign test $p = 1.9 \cdot 10^{-6}$, $p = 1.3 \cdot 10^{-4}$ after Bonferroni correction over 66 domains), justifying the targeted per-domain policy structure.

4.3 Mechanism Analysis - Decoupling TR Size from Search Cost

By instrumenting our symbolic planner to log per-instance BDD node statistics at preprocessing and throughout search, we reveal the core mechanism governing these results:

1. **Successful Design Target:** The ordering successfully achieves its explicit design objective: Transition Relation (TR) node sizes shrink significantly (geometric mean ratio 0.867, median 0.954, reduced on 58.7% of the 1,361 active instances with a comparable pair, Wilcoxon signed-rank $p \approx 3 \cdot 10^{-21}$).
2. **Invariance of Operative Bottlenecks:** The ordering leaves peak search-BDD node sizes unchanged (geometric mean ratio 1.026, median 1.000, $p = 0.12$, $n = 759$).
3. **Decoupling of TR Size and Search Cost:** TR node size is a poor proxy for runtime, whereas peak search-BDD size is the true operative bottleneck. The correlation between $\log$ TR size ratio and $\log$ runtime ratio is weak ($\rho = 0.202$, $R^2 \approx 0.04$), whereas peak search-BDD node ratio strongly predicts runtime ratio ($\rho = 0.593$, $p \approx 2 \cdot 10^{-69}$; both computed on

the 718 active instances whose TR size changed, and 0.199 / 0.580 on all 759; Figure 2 plots the latter).

4. Cost Asymmetry: On 36.7% of active instances where the ordering increases TR size (median growth 1.20 $\times$), TR expansion incurs a higher runtime penalty (time ratio 1.211) than the speedup gained from TR compression (time ratio 0.946).
5. Functional Scope: The method is a structural no-op on 468 of the 1,830 instances that report an edge count (25.6%; 13 of 66 domains) due to the complete absence of cross-variable mutex groups.

4.4 Empirical Ceiling Bounds & Portfolio Selection Limits

We investigated whether algorithm selection, restart schedules, or portfolio wrappers could extract suite-wide gains from constraint-aware variable orderings. We established several theoretical and empirical ceiling bounds:

- Run-to-run variability: The selector executes the stock code path whenever it declines to switch, giving 1,076 pairs of runs of *identical* code; their run-times have $\text{sd}(\log \text{ratio}) = 0.231$, and 33% differ by more than 20%. We emphasise that this measures *timing* dispersion only. We do not have a clean full-suite replicate of a single configuration, and therefore cannot estimate a coverage-resolution floor from our data; the corresponding caution is empirical rather than derived. It is not idle: an earlier 551-instance sample suggested a +3 suite-wide gain that did not replicate at full scale. We therefore treat small suite-wide coverage differences as uninformative without a replicate to calibrate them, and we do not claim a numerical resolution bound.
- Oracle Bound: An ideal 2-ordering oracle selector – one that picks the better of stock and constraint-aware ordering per instance, at zero cost – yields +13 instances (1,158 vs. 1,145), i.e. 0.7% of the suite. Extrapolating from the measured per-instance dispersion, a k -ordering oracle is *estimated* to reach roughly +20 to +33 instances for $k = 3 \dots 8$. These projections are model estimates, not measurements, and we did not release the simulator that produced them; they should be read as indicative only. What the released data do support is narrower: the achievable ceiling is a fraction of a percent of the suite, and the realised selector (below) captures none of it.
- Restart Schedules: In sequential budget-splitting tests under an 1800 s limit, only one stock-unsolved instance is rescued by the constraint ordering within 600 s (four within 900 s), while 50 instances that stock solves require > 900 s. Consequently, all tested budget splits reduce overall coverage.
- TR-Probe Selection: A build-time probe selector evaluating initial TR node size selects the faster ordering in 43 of the 65 deviating picks for which both arms independently solve (66%; unclustered sign test $p = 0.013$, but $p \approx 0.08$ once resolved by domain). Given that the rule’s threshold was itself tuned on these data, we treat this as exploratory rather than an established signal. Measured end-to-end, however, the wrapper is **18.4% slower** than stock in wall-clock (Table 3) while its search-only ratio is 0.99. The deficit therefore lies outside measured search time, but our data cannot apportion it among probing, process startup, repeated preprocessing and caching: even the 759 instances solved during pre-solve, which run no probe at all, show a wall ratio of 1.16 against a search ratio of 1.01. Its nominal +5 coverage (1150 vs. 1145) is within plausible run-to-run variation and is not a claim we make.

4.5 Methodological Findings

Our investigation revealed two methodological insights for automated planning evaluations:

4.5.1 Demonstration of Regression to the Mean

Selecting “hard” instances based on baseline runtime artificially skews comparative speedup claims. Filtering to instances where stock SymK takes > 300 s makes constraint ordering appear 20% faster (wall ratio 0.803, $n = 90$). Conversely, filtering by the constraint ordering’s *own* runtime makes the same method appear 10% slower (1.096, $n = 94$). A symmetric filter – both configurations above the threshold – yields 0.919 ($n = 80$, faster on 54%), and the effect shrinks toward unity as the threshold falls (0.964 vs. 1.025 at > 10 s). Standard asymmetric “hard instance” subsetting produces pure regression-to-the-mean artifacts.

4.5.2 IPC Cross-Year Duplicate Leakage

IPC benchmark instances recur across competition years. Of the 1,847 instances we evaluate, **322 are byte-identical to an instance filed under a different domain** (155 duplicate groups, by MD5 over the problem files). Leave-one-domain-out protocols therefore train and test on literally the same tasks, and per-domain results are not fully independent observations; deduplication should precede any cross-validated claim on this suite. For our own data the hazard is moot because there is little to predict in the first place: ridge regression from preprocessing features (variable count, co-occurrence edge count and density, mutex-BDD and transition-relation sizes) to the ordering’s log runtime effect reaches only $R^2 = 0.06$ under leave-one-domain-out and $R^2 = 0.05$ under leave-one-*family*-out cross-validation, consistent with Kissmann and Hoffmann’s verdict [5] that ordering quality is not predictable before the BDDs are built.




Figure 2: **Why the ordering does not pay off: it moves the wrong quantity.** Each point is one instance on which the method is active and both configurations solve ($n = 759$); axes are log ratios of constraint-aware ordering to the GAMER baseline, with the least-squares trend in black. **Left:** the transition-relation size the objective actually optimizes barely predicts runtime ($\rho = 0.199$) – a diffuse cloud around a nearly flat trend. **Right:** peak search-BDD size, which the ordering leaves essentially unchanged, predicts runtime strongly ($\rho = 0.580$). Shrinking the transition relation is not shrinking the search.

Table 1: Pre-registered confirmatory A/B evaluation results (480 total runs across 2 execution replicates).

Pre-Registered Gate	Target Requirement	Empirical Outcome	Status
Primary: Coverage Contrast	Coverage Gain $\geq +4$	+6 (Rep 1) / +6 (Rep 2)	CONFIRMED
Secondary: Woodworking Speed	Wall Time Geomean ≤ 0.85	0.647 ($1.55\times$ speedup, $n = 94$)	CONFIRMED
Guard: Domain Losses	No Replicated Domain Loss	0 Replicated Domain Losses	PASSED
Integrity: Cost Agreement	No cost disagreement	0 / 129 comparable pairs	PASSED

5. Discussion & Related Work

5.1 Symbolic Search and BDD Variable Ordering

Symbolic state-space exploration relies on Binary Decision Diagrams (BDDs) to represent state sets and transition relations compactly. The computational cost of symbolic search is dictated by the BDD node count, which is notoriously sensitive to the chosen variable ordering. GAMER [3] introduced local search over a weighted linear arrangement objective on causal graph edges to optimize BDD variable orders. This scheme was inherited largely unchanged by state-of-the-art planners, including SymBA* [10] and SymK [7, 6].

Kissmann and Hoffmann [5] conducted an extensive study of BDD variable orderings in classical planning, establishing that ordering quality can vary wildly across instances and that finding optimal orderings statically is highly unpredictable. Our results sharpen their verdict by identifying *which* quantity resists prediction: not BDD size in general, but specifically the peak size reached during search, which no preprocessing-time property we measured anticipates (Section 4.3).

5.2 Invariants and Constraint-Aware Ordering

The inclusion of state invariants, such as h^2 mutexes and exactly-one groups, is essential for pruning spurious states during symbolic search frontier expansion [9]. Torralba and Alcázar [9] observed that the constraint BDDs depend on the variable ordering and suggested exploring orders derived from the constraints. We give a concrete instantiation of that direction and evaluate it inside SymK.

Outside classical planning, constraint-graph variable ordering is well-established in SAT and VLSI CAD literature, exemplified by force-directed heuristics like FORCE and MINCE [1]. While approaches like GamerPre [5] incorporate precondition co-occurrence edges, and Torralba et al. [11] use mutexes and invariants for pruning and constraint-BDD compilation while holding the variable order fixed, our framework specifically targets state-invariant mutex and invariant group co-occurrences within the GAMER arrangement objective itself.

5.3 Algorithm Selection, Portfolios, and Methodology

Per-instance algorithm selection and sequential portfolios (e.g., SATzilla [12], Fast Downward Stone Soup [4]) are standard techniques for navigating solver complementary strengths. Both presuppose components with comple-

Table 2: Per-domain coverage breakdown on the five signal domains across Replicates 1 and 2. Counts are instances solved within 1800s; $|\mathcal{I}|$ is the number of instances in the domain.

Domain	Weight w_{co}	$ \mathcal{I} $	Stock (R1/R2)	Policy (R1/R2)	Coverage Δ
<i>parking-opt14-strips</i>	1.0	20	1 / 0	3 / 3	+2 / +3
<i>barman-opt11-strips</i>	1.0	20	9 / 9	10 / 11	+1 / +2
<i>tpp</i>	1.0	30	8 / 8	9 / 9	+1 / +1
<i>woodworking-opt08-strips</i>	2.0	30	27 / 28	28 / 28	+1 / +0
<i>woodworking-opt11-strips</i>	2.0	20	19 / 20	20 / 20	+1 / +0
Total Signal Domains	–	120	64 / 65	70 / 71	+6 / +6

Table 3: Suite-wide performance across all 1,847 IPC optimal-STRIPS instances. Time ratios are **wall-clock relative to stock**, including all selector probe overhead, over instances solved by both the row configuration and stock; intervals are domain-clustered bootstrap 95% CIs. The selector’s search-only ratio is 0.992: its deficit lies outside measured search time (probing, process startup, repeated preprocessing), which these data cannot further apportion.

Configuration	Coverage	Geomean Wall Ratio	Median Ratio	95% Conf. Interval	n
Stock SymK (Baseline)	1145	1.0000	1.0000	–	–
Always-On Constraint Ordering	1139	1.0053	1.0000	[0.97, 1.05]	1126
TR-Probe Selector	1150	1.1837	1.1740	[1.16, 1.21]	1141

mentary strengths and a signal to select on. Our analysis (Section 4.4) indicates that neither precondition holds strongly here: the entire oracle gain available from these two orderings is +13 instances, and the selector we built captures none of it net of its own overhead.

5.4 Limitations

Three limitations scope our claims. (i) The five signal domains were identified in the same full-suite runs used for the suite-wide analysis; the pre-registered A/B therefore establishes that the gain on *these* domains is reproducible rather than noise, not that a policy derived this way would transfer to unseen domains. (ii) All tasks ran under the IPC optimal-track limits (1800s, 8 GB, one core), but on a single shared multi-core server; since coverage at a fixed time limit also depends on hardware and load, absolute totals are not directly comparable across papers – within-run contrasts, which cancel these factors, are the quantitatively meaningful ones. (iii) The k -ordering ceiling projections are model extrapolations, not measurements.

6. Conclusion & Future Work

In this paper, we implemented and evaluated constraint-aware variable ordering inside SymK across the complete IPC optimal-STRIPS benchmark suite. Our inquiry delivers four primary outcomes:

- **Reproducible Per-Domain Gain:** A pre-registered, double-replicated A/B policy gains +6 **coverage** in both replicates on its five post-hoc-selected target domains and accelerates *woodworking* by 1.55 $\times$, with no cost disagreement. This is

same-task reproducibility on selected domains, not a general planner advance.

- **Bounded Suite-Wide Null:** Enabled universally, the ordering yields a tightly bounded zero effect (coverage 1145 vs. 1139; geomean wall ratio 1.005, 95% CI [0.97, 1.05]), excluding even a 5% suite-wide effect.
- **Mechanism Uncovering:** The ordering compresses transition relations by 13.3% as designed, but TR size explains only $\approx 4\%$ of the variance in runtime; peak search-BDD size, which the ordering does not move, is the operative quantity. Optimizing the wrong proxy is why a sound idea produces no effect.
- **Ceiling Estimates & Methodological Insights:** The oracle over both orderings is worth +13 instances (0.7% of the suite), and our probe-based selector captures none of it net of overhead – so selection wrappers are unpromising here, though we stop short of a formal impossibility bound. We also document traps in hard-instance subsetting and IPC duplicate leakage.

6.1 Future Directions

Our diagnostic decoupling points directly toward two promising avenues for future research in symbolic search:

6.1.1 Targeting Peak Search-BDD Size

Because peak search-BDD node count is highly predictive of runtime ($\rho = 0.593$) and bit-deterministic across runs ($r = 0.999$), future variable ordering objectives should target peak search-BDD expansion directly. Since peak

BDD sizes are unobservable during preprocessing, developing lightweight online probes or dynamic reordering criteria that estimate search-space width remains an open challenge.

6.1.2 Exploiting Ordering Diversity

Published random-ordering distributions exhibit heavy favorable tails that correlated structural orderings cannot reach [5]. Designing randomized, high-diversity ordering portfolios paired with early-termination rerun controls could unlock the heavy-tail performance gains of symbolic BDD search.

References

- [1] Fadi A. Aloul, Igor L. Markov, and Karem A. Sakallah. FORCE: A fast and easy-to-implement variable-ordering heuristic. In *Proceedings of the 13th ACM Great Lakes Symposium on VLSI (GLSVLSI)*, pages 116–119, 2003.
- [2] Dominik Drexler, Daniel Gnad, Paul Höft, Jendrik Seipp, David Speck, and Simon Ståhlberg. Ragnarok. In *Tenth International Planning Competition (IPC-10): Planner Abstracts, Classical Tracks*, 2023.
- [3] Stefan Edelkamp and Peter Kissmann. GAMER: Bridging planning and general game playing with symbolic search. In *6th International Planning Competition (IPC-6)*, ICAPS, 2008.
- [4] Malte Helmert, Gabriele Röger, and Erez Karpas. Fast downward stone soup: A baseline for building planner portfolios. In *ICAPS 2011 Workshop on Planning and Learning (PAL)*, pages 28–35, 2011.
- [5] Peter Kissmann and Jörg Hoffmann. BDD ordering heuristics for classical planning. *Journal of Artificial Intelligence Research (JAIR)*, 51:779–804, 2014.
- [6] David Speck. SymK – a versatile symbolic search planner. In *Tenth International Planning Competition (IPC-10): Planner Abstracts, Classical Tracks*, 2023.
- [7] David Speck, Robert Mattmüller, and Bernhard Nebel. Symbolic top-k planning. In *Proceedings of the Thirty-Fourth AAAI Conference on Artificial Intelligence (AAAI-20)*, pages 9967–9974, 2020.
- [8] Ayal Taitler, Ron Alford, Joan Espasa, Gregor Behnke, Daniel Fišer, Michael Gimelfarb, Florian Pommerening, Scott Sanner, Enrico Scala, Dominik Schreiber, Javier Segovia-Aguas, and Jendrik Seipp. The 2023 international planning competition. *AI Magazine*, 45(2):280–296, 2024.
- [9] Álvaro Torralba and Vidal Alcázar. Constrained symbolic search: On mutexes, BDD minimization and more. In *Proceedings of the Sixth Annual Symposium on Combinatorial Search (SoCS)*, pages 175–183, 2013.
- [10] Álvaro Torralba, Vidal Alcázar, Daniel Borrajo, Peter Kissmann, and Stefan Edelkamp. SymBA*: A symbolic bidirectional A* planner. In *Eighth International Planning Competition (IPC-8): Planner Abstracts*, pages 105–109, 2014.
- [11] Álvaro Torralba, Vidal Alcázar, Peter Kissmann, and Stefan Edelkamp. Efficient symbolic search for cost-optimal planning. *Artificial Intelligence*, 242:52–79, 2017.
- [12] Lin Xu, Frank Hutter, Holger H. Hoos, and Kevin Leyton-Brown. SATzilla: Portfolio-based algorithm selection for SAT. *Journal of Artificial Intelligence Research (JAIR)*, 32:565–606, 2008.